
CineIQ: Enterprise-Grade Hybrid Recommendation Engine

A multi stage hybrid pipeline with contextual sentiment reranking and explainable visual analytics

Technical Project Report

Prepared by:

Arnav De (Roll No: 240103022)

Arya Sandeep Kshirsagar (Roll No: 240103023)

Coding Club, IIT-G

May 18, 2026

1 Introduction

CineIQ is a multi stage movie recommendation system engineered to seamlessly synthesize behavioral learning, high dimensional metadata similarity vectors, and transparent, explainable AI analytics within a singular, unified pipeline. Rather than relying on an isolated algorithmic paradigm, the architecture integrates a decoupled, multi layered execution flow consisting of **collaborative filtering via Singular Value Decomposition (SVD)**, **content-based filtering utilizing Term Frequency-Inverse Document Frequency (TF-IDF) text vectorization with linear kernel similarity**, a data-driven **weighted hybrid stacking ensemble**, and an advanced **transformer-driven sentiment re-ranking engine**. To bridge the gap between machine optimization and human interpretability, the pipeline deploys a production-tier **Streamlit analytics dashboard** that maps historical user taste profiles and routes dynamic, rule-based matching explanations in human-readable formats.

The pipeline was validated locally utilizing the **MovieLens 25M** interaction database, comprehensive **TMDB movie metadata** for semantic text mining, and the **IMDB 50K text corpus** for sentiment benchmarking. In the finalized implementation, the data engineering layer successfully processed, downcasted, and serialized **25,000,095 raw rating interactions** in exactly **65.89 seconds**. The collaborative modeling track isolated a controlled **20% randomized sample comprising 5,000,019 interactions** for localized hyperparameter tuning, where the trained SVD model achieved a highly stable validation **RMSE of 0.8758**. These empirical milestones confirm that CineIQ functions as a highly optimized, fully validated pipeline capable of handling data processing and inference tasks at scale.

2 Datasets and Technology Stack

The engine ingests data from three distinct sources, with each source engineered to feed a dedicated abstraction layer of the hybrid recommendation framework:

1. **MovieLens 25M Database:** The primary behavioral interaction matrix, encapsulating approximately **25,000,095 user-movie numerical ratings**. This dense matrix serves as the fundamental foundation for latent factor decomposition, user preference profiling, and collaborative filtering.
2. **TMDB Metadata Corpus:** Provides rich, movie-level attribute profiles containing dense structural features such as **genres, cast configurations, production crews, descriptive plot keywords, release dates, and financial performance (revenue)**. These multi-modal attributes are leveraged for high-dimensional semantic feature extraction and front-end dashboard visualization.
3. **IMDB 50K Reviews Matrix:** Comprises **50,000 highly polar, text-heavy cinematic reviews**. This dataset is isolated exclusively as an objective validation benchmark to audit lexicon-based and contextual deep-learning sentiment models prior to pipeline deployment.

The infrastructure was constructed natively in **Python** and is segmented into specialized functional domains to maintain a clean, decoupled development lifecycle:

- **Data Engineering & Storage:** Pandas, NumPy, PyArrow, and FastParquet are utilized to execute aggressive numerical datatype downcasting and handle dense, columnar binary serialization.
- **Statistical Optimization & Vectorization:** Scikit-Learn handles sparse high-dimensional TF-IDF vector text mining and fits the Ordinary Least Squares (OLS) Linear Regression meta-model.

- **Collaborative Modeling:** The **Surprise** framework is deployed to handle sparse matrix configurations and compute optimized SVD weights.
- **Natural Language Processing:** **NLTK VADER** and **Hugging Face Transformers (DistilBERT)** orchestrate the lexical and contextual sentiment inference routines.
- **MLOps Tracking & Execution UI:** **MLflow** acts as the persistent experiment tracking server backed by a local relational database engine, while **Streamlit** and **Plotly** provide the interactive, explainable visual frontend.




Figure 1: CineIQ: Complete Recommendation Pipeline Workflow

3 System Design and Implementation

CineIQ is architected as a decoupled, multi-staged machine learning pipeline, where each operational phase is isolated into a dedicated script within the repository. This modular separation enforces strict boundary controls, mitigates memory allocation spikes, and establishes clean checkpoints for localized experiment tracking.

3.1 Phase 0: Data Engineering and Preprocessing

The data engineering layer is responsible for translating massive, multi-modal raw files into highly optimized, localized view states. Because loading 25 million string and float entries simultaneously exceeds memory thresholds on standard consumer hardware, the module executes strict numerical downcasting—converting standard 64-bit float and integer fields into compact, memory-efficient 32-bit blocks.

This phase merges the raw behavioral logs with textual metadata to execute two core transformations: tracking the raw structural matrix shapes and mining semantic strings. To build the content matching space, the script aggregates space-stripped and lowercased categorical fields (directors, actors, genres, keywords) and appends them to the free-text movie plot summary, synthesizing a unified text “soup.” To prevent redundant runtime parsing, the data engineering block serializes and dumps four static, compressed binary structures directly into cold storage:

- `svd_view.parquet`: A downcasted three-column matrix preserving the full **25,000,095-row** behavioral interaction history.
- `content_view.parquet`: Mapping vector associating unique movie IDs with their parsed textual soup strings.
- `sentiment_view.parquet`: A cleaned evaluation matrix containing pre-mapped binary sentiment targets (1 for positive, 0 for negative).
- `dashboard_view.parquet`: A high-speed analytical lookup frame storing presentation meta-data like titles, release dates, and financial revenues.

By utilizing memory-mapped columnar formatting, the entire processing operation over the 25M row input matrix completes execution in an optimized duration of exactly **65.89 seconds**.

3.2 Phase 1A: Behavioral Collaborative Filtering

The collaborative filtering track captures latent behavioral dynamics between user profiles and movie items using Singular Value Decomposition (SVD). To ensure training remains highly performant and feasible within localized memory allocations, the model trains on a seed-controlled 20% randomized sample representing **5,000,019 interaction cells**.

The optimization routine factorizes the sparse target rating matrix into low-rank user and item embeddings by minimizing a regularized mean squared error loss. The low-rank approximation maps a rating estimate $\hat{r}_{u,i}$ via the scalar dot product:

$$\hat{r}_{u,i} = \mu + b_u + b_i + P_u \cdot Q_i^T \quad (1)$$

Where μ is the global baseline rating, b_u and b_i represent user and item biases, and P_u and Q_i represent the regularized latent factor matrices mapped across $n_factors = 100$. The model achieved a stable, verified validation **RMSE of 0.8758**.

To log performance metrics, the module integrates an active **MLflow** tracking connection. It records hyperparameters ($n_factors$, $random_state$) and splits timings into two distinct logged telemetry metrics: `cv_training_time_seconds` and `production_retrain_time_seconds`. Following production retraining, the script evaluates candidate rating predictions across a designated target block of **500 active users** and caches the top 100 unrated candidates per user in an MLflow versioned artifact at `collaborative_scores.pkl`.

3.3 Phase 1B: Semantic Content-Based Filtering

Operating independently from the behavioral matrix, the content-based engine maps semantic alignment by analyzing text similarity spaces. The corpus vocabulary space is bounded using a maximum feature constraint ($max_features = 5000$) running over a scikit-learn `TfidfVectorizer` initialized with English stop-word filtering.

To build a contextual profile for an active user, the script isolates the user’s historical interaction rows to pull out favorable ratings (≥ 3.5). It queries the corresponding movie rows from the global TF-IDF sparse array, scales each matrix line by its true numerical rating value, and computes a unified **User Taste Vector**:

$$\vec{U}_{\text{taste}} = \frac{\sum_{i \in R_+} r_i \cdot \vec{M}_i}{\sum_{i \in R_+} r_i} \quad (2)$$

Where R_+ represents the set of favorably rated movies and r_i denotes the explicit interaction score. High-speed vector similarity scores are computed across the unrated movie space by taking the linear kernel dot product:

$$\text{Similarity}(\vec{U}_{\text{taste}}, \vec{M}_{\text{candidate}}) = \vec{U}_{\text{taste}} \cdot \vec{M}_{\text{candidate}}^T \quad (3)$$

The vectorizer configuration parameter and the evaluated matrix dimensions are registered as parameters in `MLflow`, while the final top 100 candidate items per user are written out to `content_scores.pkl`.

3.4 Phase 2: Data-Driven Stacking Ensemble

To merge behavioral factor ratings with high-dimensional text similarities without relying on arbitrary human guesswork, CineIQ implements an objective, data-driven stacking fusion methodology. The script generates an evaluation meta-dataset containing 50,000 historical sample rows, matching actual historical ratings against the prediction vectors output by the independent SVD and TF-IDF models.

An Ordinary Least Squares (OLS) Linear Regression meta-model is fitted against this feature space to automatically learn the true statistical significance of both indicators. The empirical optimization cycle resolved the following tracking policy:

$$S_{\text{Hybrid}} = (0.9578 \cdot S_{\text{Collaborative}}) + (4.4236 \cdot S_{\text{Content}}) - 0.2676 \quad (4)$$

The higher weight assigned to the content score acts as a mathematical normalization adjustment. High-dimensional sparse TF-IDF cosine products naturally generate smaller raw variations compared to dense SVD metrics. The meta-model uncovers this scaling variance and recalibrates the weights automatically to ensure both recommendation tracks contribute equitably to the final hybrid score payload (`hybrid_scores.pkl`).

3.5 Phase 3: Contextual Transformer Sentiment Re-Ranking

The final algorithmic refinement stage applies an additional quality adjustment layer. Before processing, a sentiment auditing benchmark was conducted on the 50,000-row IMDB corpus to contrast lexicon methods against state-of-the-art transformer contextual engines. The empirical results demonstrate a clear architectural advantage:

Sentiment Engine Architecture	Validation Accuracy
Lexical VADER Engine (CPU Baseline)	69.70%
DistilBERT Pipeline (GPU Baseline)	88.25%

During pipeline processing, the script extracts the unique union-set of movie candidate IDs recommended across the user profiles (shrinking the payload from 50,000 entries down to a highly optimized block of 3,554 unique entities to protect system memory limits). It passes the metadata soups through the pipeline to compute a scaled polarity metric ($s_i \in [-1, 1]$). The final ranked recommendation score is computed using a dynamic catalyst adjustment parameter ($\alpha = 0.1$):

$$S_{\text{Final}} = S_{\text{Hybrid}} \cdot (1.0 + (\alpha \cdot s_i)) \quad (5)$$

The runtime parameters, validation accuracy logs, and inference times are piped to the local SQLite serverless backend tracker, and the finished dataset is registered as a production asset at `final_ranked_scores.pkl`.

3.6 Phase 4: Explainable Dashboard Architecture

The user interface layer provides complete visibility into the underlying model data. Built with Streamlit, the dashboard visualizes a user’s historical profile—displaying genre footprint visualizations, decade preferences, and top director charts—alongside ranked recommendation cards.

To build trust, the dashboard utilizes a rule-based routing engine that displays a transparent, human-readable explanation card for each item based on its internal scores:

- **Strong Behavioral Match:** Highlighted when the collaborative SVD prediction profile dominates the metrics.
- **Semantic Taste Alignment:** Triggered when text similarities reveal heavy overlap with historical preferences.
- **Consensus Fit:** Displayed when both model tracks yield strong matching probabilities.
- **Public Acclaim Catalyst:** Triggered when a high candidate placement is driven by excellent DistilBERT sentiment scores.

To ensure stability under tight hardware ceilings, the frontend uses caching decorators (`@st.cache_data`) and forces garbage collection (`gc.collect()`) immediately after rendering heavy visualizations to prevent out-of-memory crashes.

4 Key Outcomes and System Performance

The deployment and empirical evaluation of the CineIQ architecture yielded several measurable, high-performance outcomes across the data engineering and algorithmic modeling tracks. By shifting from standalone scripts to an integrated, checkpointed infrastructure, the project established clear technical milestones:

Table 1: CineIQ Core Performance and Data Engineering Metrics

Pipeline Phase	Target Evaluation Metric	Verified System Output
Phase 0	Total Ingested Rating Interactions	25,000,095 rows
Phase 0	Optimization I/O Speed (Downcasting)	65.89 seconds
Phase 1A	Sampled SVD Training Subset Space	5,000,019 rows
Phase 1A	Baseline Factorization Performance	0.8758 RMSE
Phase 1A	Batch Precomputation Target Scope	500 active users
Phase 1B	Text Vectorizer Sparse Vocabulary Space	5,000 maximum features
Phase 2	Meta-Model Linear Regression Intercept	-0.2676
Phase 3	DistilBERT Contextual Classification Accuracy	88.25%

These system configurations and empirical indicators demonstrate that CineIQ successfully moves past purely isolated model calculations. By linking each block—from memory-optimized data views to regularized matrix factorization, linear stacking optimization, contextual NLP catalyst reranking, and cache-aggregated UI loops—the framework functions as a reliable,

production-ready recommender system. Furthermore, every metric and serialized model checkpoint displayed in Table 1 is permanently registered inside the local **MLflow** tracking server, providing complete version control and run reproducibility for the entire asset lifecycle.

5 Conclusion

CineIQ demonstrates the successful design, optimization, and deployment of a multi-stage hybrid movie recommendation system engineered to run efficiently within constrained, localized hardware environments. By replacing fixed heuristic rules with a decoupled, modular architecture, the system successfully unifies large-scale behavioral dynamics, high-dimensional metadata similarity spaces, and deep transformer sentiment auditing into a single, automated execution flow. The primary contributions of this work include its **memory-mapped multi-stage pipeline design**, its data-driven **learned ensemble stacking strategy**, its contextual **NLP catalyst re-ranking loop**, and its user-centric, **explainable visual analytics interface**.

From a software engineering perspective, the project proves that processing over **25,000,095 raw interactions** and training a regularized collaborative model on a **5,000,019-row sampled matrix** is entirely feasible on local consumer hardware when backed by defensive data engineering. Implementing explicit numerical downcasting, memory-efficient Parquet serialization, candidate union filtering, and forced garbage collection successfully eliminated runtime out-of-memory faults. Furthermore, incorporating objective baseline benchmarking (achieving **88.25% classification accuracy via DistilBERT**) and transparent, rule-based explanation templates extends the system well beyond traditional black-box filtering approaches.

In conclusion, CineIQ serves as a robust, highly scalable recommendation engine prototype that bridges the gap between raw machine optimization and human interpretability. By registering all hyperparameter variants, execution durations, and serialization assets within a persistent **MLflow** tracking framework, the pipeline establishes a reliable and completely reproducible foundation. This architecture provides a clear path forward for future enhancements in distributed deployment, real-time streaming inference, and deep graph neural network embedding integration.